माधव प्रौद्योगिकी एवं विज्ञान संस्थान, ग्वालियर
MADHAV INSTITUTE OF TECHNOLOGY & SCIENCE, GWALIOR
Deemed to be University
(Declared under Distinct Category by Ministry of Education, Government of India)
NAAC ACCREDITED WITH A++ Grade
Gola Ka Mandir, Gwalior (M.P.)- 474005, INDIA
Ph.:+91-751-2409300, E-mail: vicechancellor@mitsgwalior.in, Website: www.mitsgwalior.in

Confidence-Aware RAG Pipeline with Hallucination Detection

Answer only when the sources prove it, and say why when they don't.

Shivam Gupta (BTCD24O1060) · Cornerstone Project · Madhav Institute of Technology & Science, Gwalior

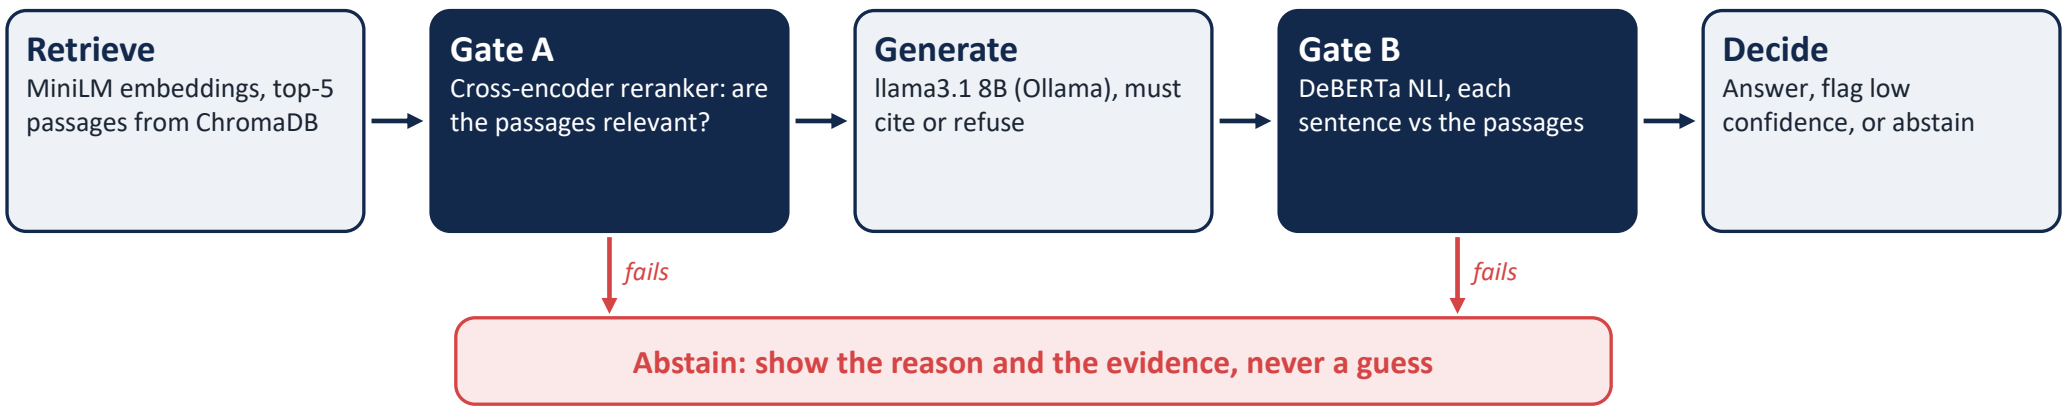
The problem

Retrieval-augmented chatbots still hallucinate. They answer in the same confident tone whether or not the retrieved passages support the answer, so three failures reach users: questions the documents cannot answer, invented details, and statements that contradict the source. A confident wrong answer costs more than an honest "I don't know".

Objectives

- Check passage relevance before generating (Gate A)
- Verify each answer sentence against the sources (Gate B)
- Abstain with a stated reason when confidence is low
- Compare against plain RAG on SQuAD 2.0 and HotpotQA

Methodology: two gates, one decision



Combined confidence is the weighted harmonic mean of Gate A (0.4) and Gate B (0.6): one weak gate drags it down, so good retrieval cannot hide an unsupported answer. A contradicted sentence caps Gate B at 0.15.

Key features

Refuses before guessing

Out-of-scope questions stop at Gate A, before the language model ever runs.

Sentence-level evidence

Every sentence labelled supported, baseless or contradicted (RAGTruth).

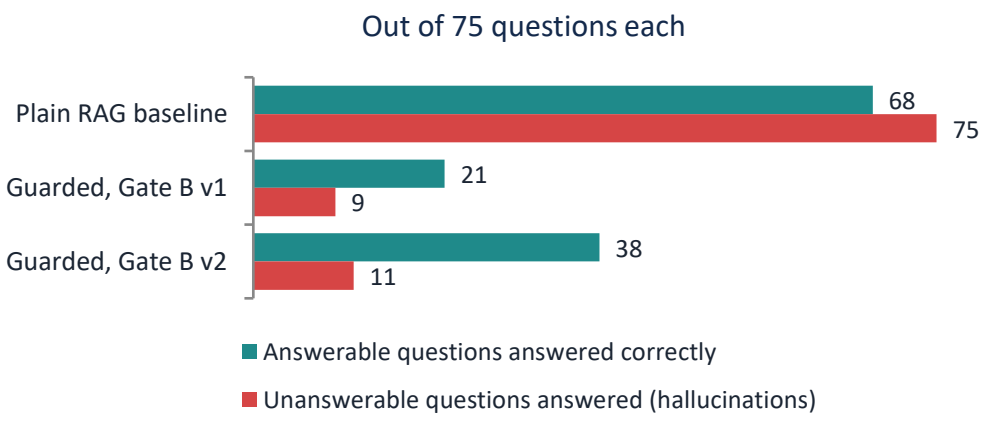
Gate B v2

Claims checked against passages and 3-sentence windows: fewer false rejections.

App and monitoring

Streamlit app shows every gate decision; a dashboard tracks trends and latency.

Results: 150 SQuAD 2.0 questions, fresh end-to-end runs



75 → 11
hallucinated answers (of 75): plain RAG vs guarded

97%
of its answers to answerable questions are correct (38 of 39)

71% → 48%
over-abstention after Gate B v2 (paired test, $p < 0.001$)

31% → 66%
HotpotQA full supporting-fact recall with chunking

Applications

- Course and campus assistants
- Policy and manual Q&A
- Healthcare, legal, finance lookups

Future scope

- Stronger NLI model (DeBERTa-large)
- Prompt-only control arm
- Multi-passage entailment for multi-hop

Built with

- Python · LangChain · ChromaDB
- DeBERTa-v3 NLI · MiniLM reranker
- Ollama llama3.1 8B · Streamlit